

Proof

Eduardo¹, Fernando¹

¹ Faculty of Informatics, Università della Svizzera italiana, Switzerland

Abstract

-

Report Info

Published

December 2012

Number

USI-INF-TR-2010-4

Institution

Faculty of Informatics

Università della Svizzera italiana

Lugano, Switzerland

Online Access

www.inf.usi.ch/techreports

1 ...

Variables are distributed among partitions based on a function $h(i)$ that maps each variable v_i to one partition $P_{h(i)}$. Therefore, no two partitions contain the same variable. The set \mathcal{C} contains all commands allowed by the application. Every command C in \mathcal{C} is associated with a variable set \mathcal{V} and a partition set \mathcal{P} . Set \mathcal{V} contains the variables accessed (read/written) by C , while \mathcal{P} is the set of partitions that hold variables in \mathcal{V} , i.e., $\mathcal{P} = \{P_{h(i)} : v_i \in \mathcal{V}\}$. C is executed by partition P_x , where x is the lowest subscript in \mathcal{P} , i.e., $x = \min_{P_i \in \mathcal{P}}(i)$.

Communication between partitions is done with oam-cast, mapping each partition to a multicast group. Such multicast protocol is quasi-genuine, requiring the creating of a *sendersTo* relation. We define it as follows: if the application allows a command C to be associated with a certain partition set \mathcal{P} , then every partition in \mathcal{P} can send to any partition in \mathcal{P} . More formally, we have:

$$\exists C \in \mathcal{C} \Rightarrow P_i \in \text{sendersTo}(P_j) : P_i \in \mathcal{P} \wedge P_j \in \mathcal{P}$$

(I'm not sure the following is necessary, but here it goes)

When a client issues a command C , it simply sends C to some replica in P_x , which will then proceed with oam-cast.

1.1 Proof

Proposition 1. *Let \mathcal{E} be any execution of SSMR. Let C_1 and C_2 be any two commands in \mathcal{E} such that $C_1 \prec C_2$ in real-time. C_2 will not precede C_1 in \mathcal{E} .*

Proof. There are x situations in which a command has to precede other in an execution. They are:

- i) C_1 and C_2 access the same variable v_x .

In this case, we have two possibilities:

- a) C_1 and C_2 are executed in the same partition.
Linearizability in this case is obvious, since there is no way to
- b) C_1 and C_2 are executed in different partitions.
In this case, necessarily there is a variable $v_y \neq v_x$, accessed by one of the two commands.
- ii) C_1 and C_2 are such that $\mathcal{P}_1 \cap \mathcal{P}_2 = \emptyset$

□